

Honeycomb: Secure and Efficient GPU Executions via Static Validation

Haohui Mai Jiacheng Zhao **Hongren Zheng[†]** Yiyang Zhao Zibin Liu
Mingyu Gao Cong Wang Huiming Cui Xiaobing Feng Christos Kozyrakis

PrivacyCore

ICT,CAS

Tsinghua[†]

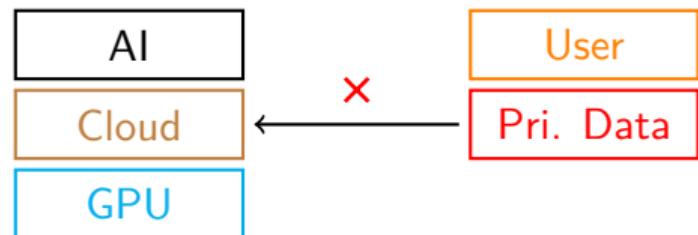
Stanford

IDEA

BUPT

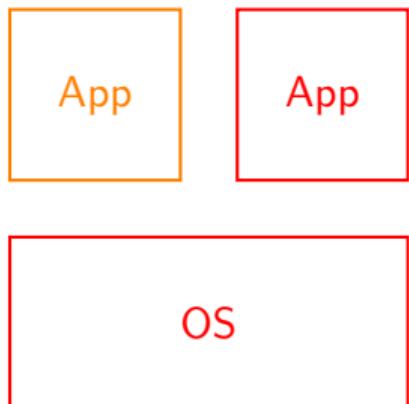
AI on private data needs S&P solutions

- AI is powerful
 - e.g. ChatGPT
- Still **security concerns**
 - Private data: e.g. medical/financial records
 - **User** does not trust 3rd party **cloud**



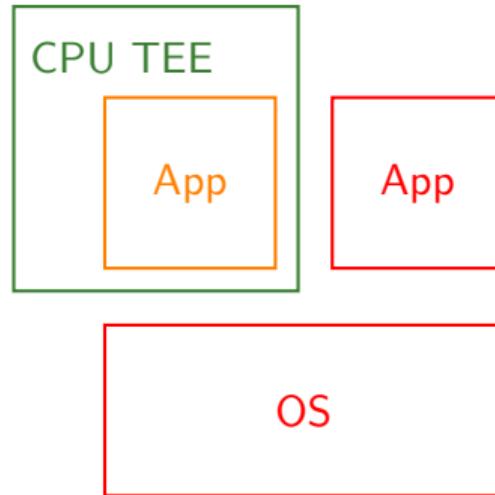
GPU TEE: a pragmatic approach

- In Cloud, User App can be harmed by other Apps and the OS



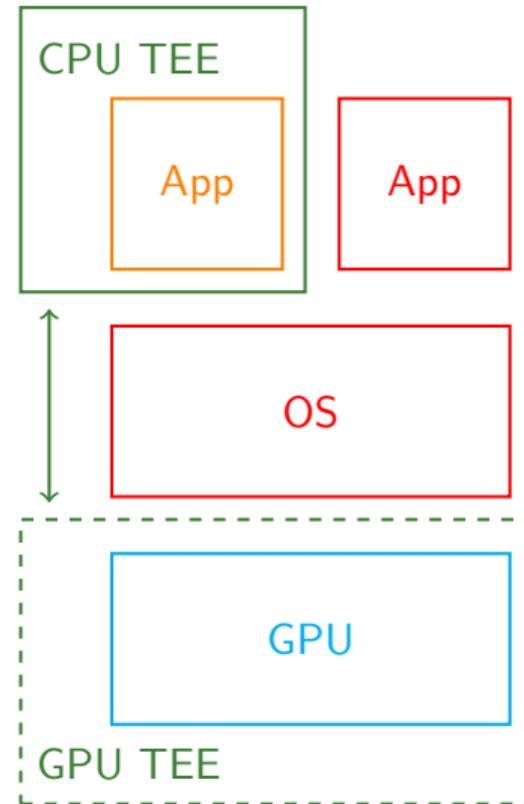
GPU TEE: a pragmatic approach

- In Cloud, User App can be harmed by other Apps and the OS
- Trusted Execution Environment (TEE) provides *isolation*
 - Special CPU hardware: Intel SGX/TDX, AMD SEV
 - Efficient: native speeds within enclave



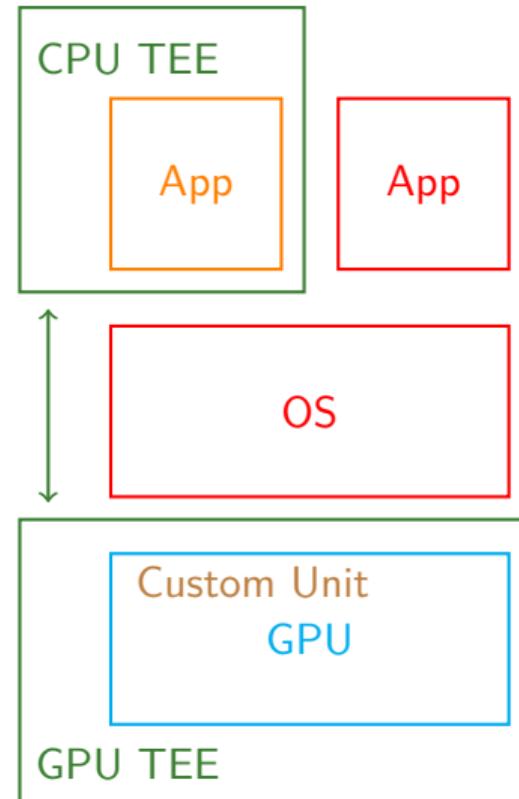
GPU TEE: a pragmatic approach

- GPU is powerful
 - Widely used by AI
- We want CPU TEE
 - Won't leak private data



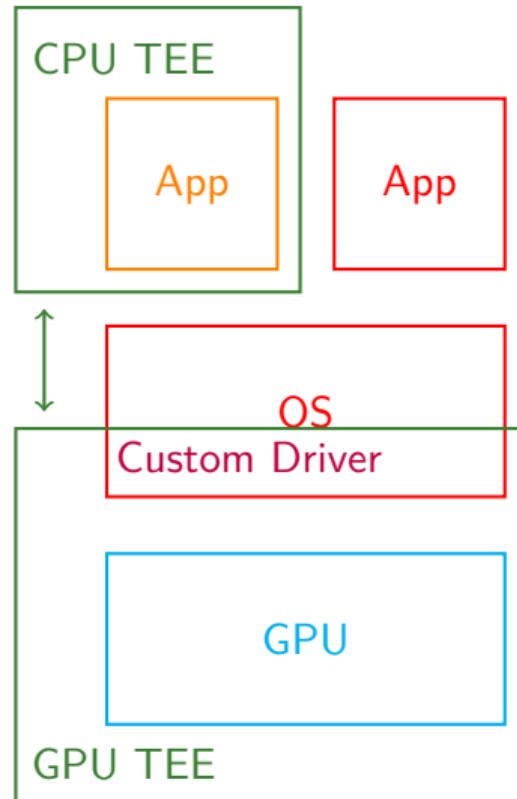
GPU TEE: a pragmatic approach

- GPU is powerful
 - Widely used by AI
- We want CPU TEE
 - Won't leak private data
- Current proposals
 - Hardware modification: slow evolution



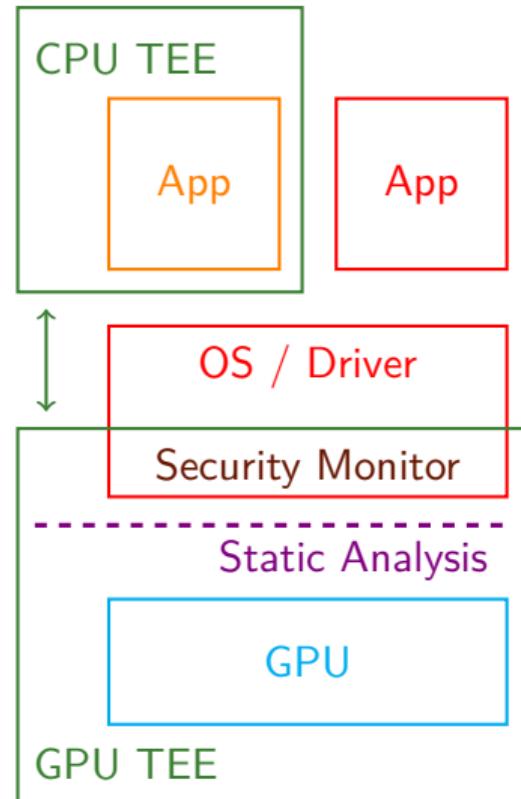
GPU TEE: a pragmatic approach

- GPU is powerful
 - Widely used by AI
- We want CPU TEE
 - Won't leak private data
- Current proposals
 - Hardware modification: slow evolution
 - Driver-based: large TCB (>1M SLOC), error-prone



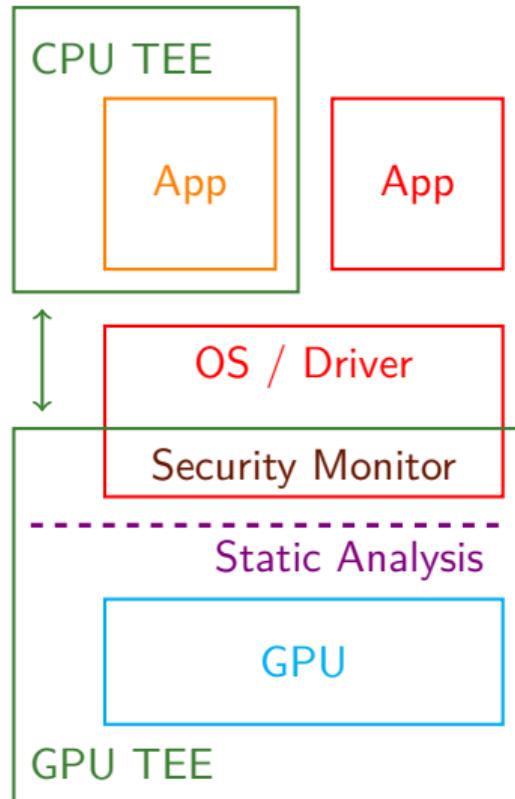
Honeycomb: confining behaviors via *static validation*

- Previous work: keep invariant
 - Either in **hardware** or **driver**
- Our work: by **static validation**
 - On the proper interface -----



Honeycomb: confining behaviors via *static validation*

- Flexible: Complement hardware limitations
- Efficient:
 - Security checks at load time: 2% overheads for BERT / NanoGPT.
 - Modest overall dev. efforts.
- Secure: 18x smaller TCB compared to Linux-based systems



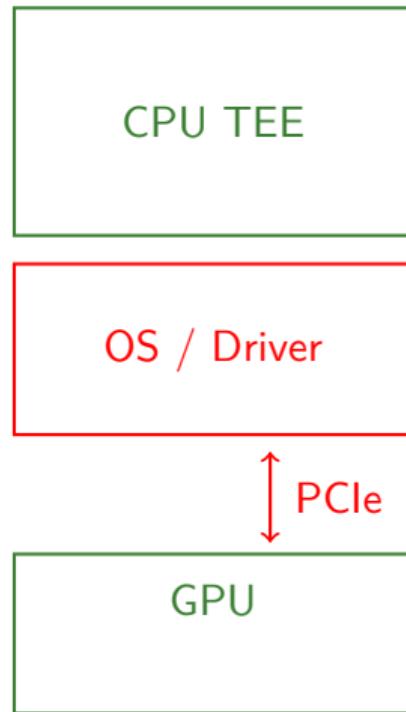
Agenda

- 1 Introduction
- 2 Assumptions & Background
- 3 Design & Implementation
- 4 Evaluation & Experience
- 5 Conclusion

Threat Model

Adversary

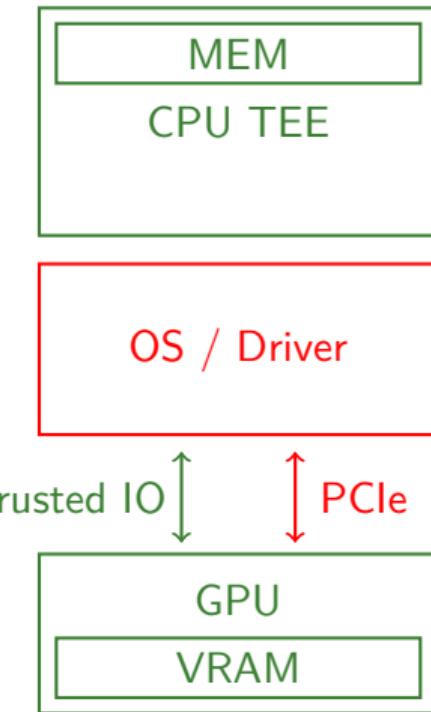
- Controls entire software stacks (OS / compiler / hypervisor)
- Has physical access of the hardware
- Sniffs PCIe traffic
- But cannot tamper the CPU or GPU silicons



Threat Model

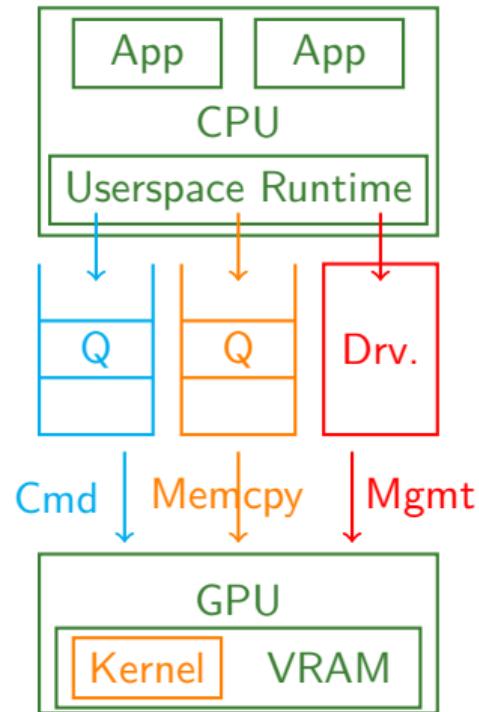
Assumptions

- CPU TEE (e.g., AMD SEV-SNP)
- Discrete GPU with integrated memory
- Trusted I/O paths: detailed in the paper
- Side-channel attacks are out of scope



GPU is a discrete accelerator

- Userspace queues
 - MemOp queue: Memcpy **kernel** / data
 - Cmd queue: Launch kernel
- Kernel space **driver**
 - Initialize hardware and address space
 - Alloc/Multiplex device memory / **queues**

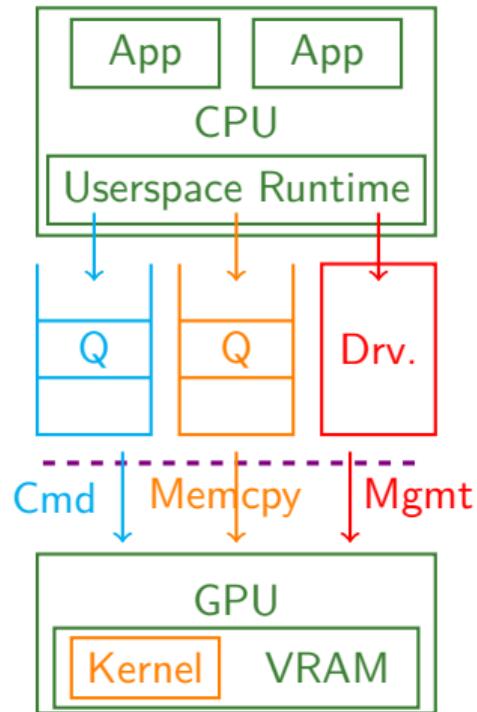


Agenda

- 1 Introduction
- 2 Assumptions & Background
- 3 Design & Implementation
- 4 Evaluation & Experience
- 5 Conclusion

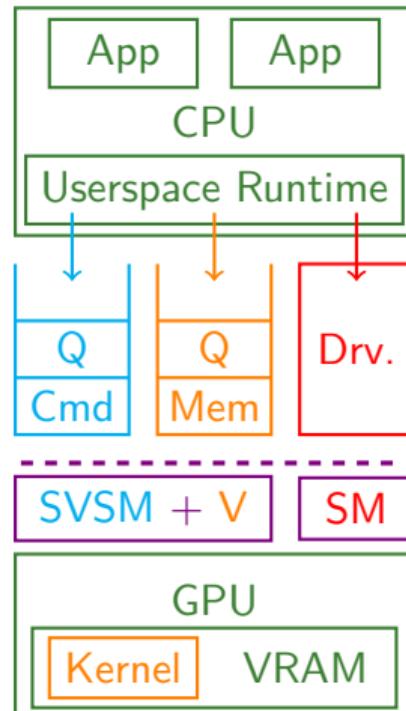
Overview of Honeycomb

- Keep invariant by **static validation**
 - On the proper interface -----
 - Regulating high-level semantics



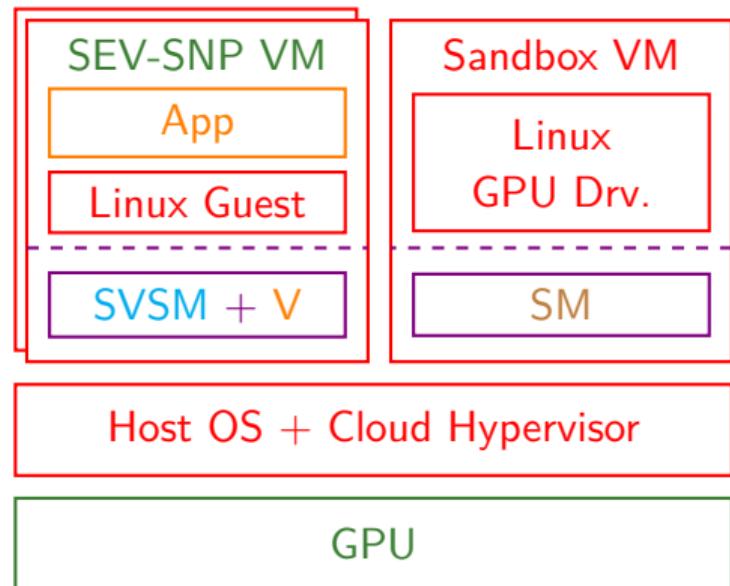
Overview of Honeycomb

- Keep invariant by **static validation**
 - On the proper interface -----
 - Regulating high-level semantics
- **SVSM**: on Queue
 - Secure VM Service Module
- **Validator**: ensure *safe* GPU kernel
- **Security Monitor**: on Driver



Architecture of Honeycomb

- Honeycomb and hardware are TCB
- SVSM for SEV-SNP VM
 - SEV-SNP is an AMD CPU TEE feature
 - Regulating user app behavior
- Validator
 - Ensure safe GPU kernels
- SM for Sandbox VM
 - Regulating GPU driver
- System-wide invariant: Efficient IPC
 - Between GPU kernels of Apps



SVSM / SM: intercepting at *lowest* level

- Validate queue cmd / MMIO
- To ensure
 - Validated kernels
 - Init sequences
 - Memory isolation
 - Secure memcpy
- Remove OS kernel / GPU driver / runtime from the TCB

```
void check_launch_kernel(  
    AddrSpace *addr,  
    DispatchPkt *p) {  
    if(!validated(addr,  
                 p->kernel_object))  
        abort_user();  
    ...
```

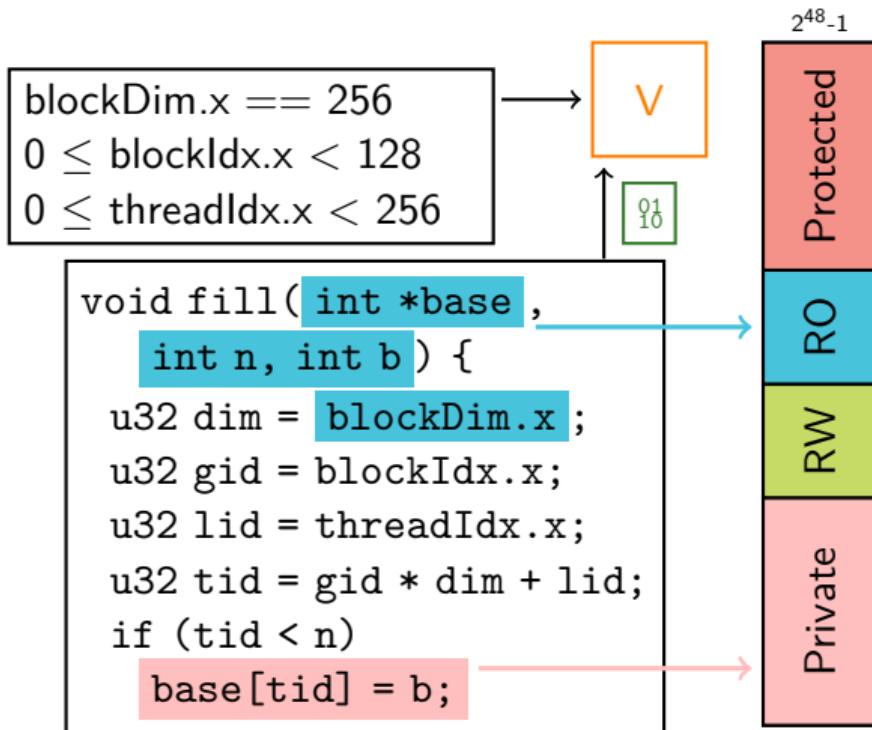
SVSM / SM: intercepting at *lowest* level

- Validate queue cmd / MMIO
- To ensure
 - Validated kernels
 - Init sequences
 - Memory isolation
 - Secure memcpy
- Remove OS kernel / GPU driver / runtime from the TCB
- Challenge: Recover semantics from MMIO ...
- BTW: Found 5 new bugs in AMDGPU, deployed in Linux 5.19

```
void check_launch_kernel(  
    AddrSpace *addr,  
    DispatchPkt *p) {  
    if(!validated(addr,  
                 p->kernel_object))  
        abort_user();  
}
```

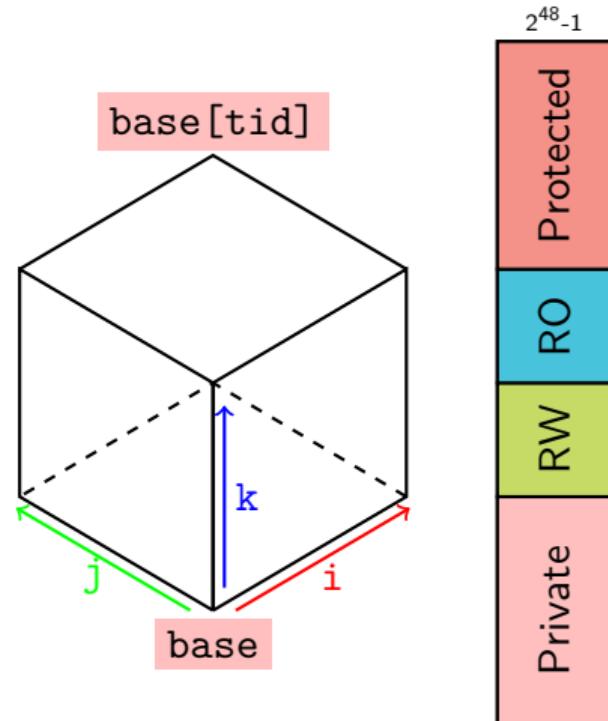
Validator: analysing *binaries*

- Invariant: partitioned addr. space
- Integrity: analyse memory access range
 - e.g. No write to protected region
- Validate GPU kernel at load time
 - Modest overhead
- On binaries: remove compiler from TCB



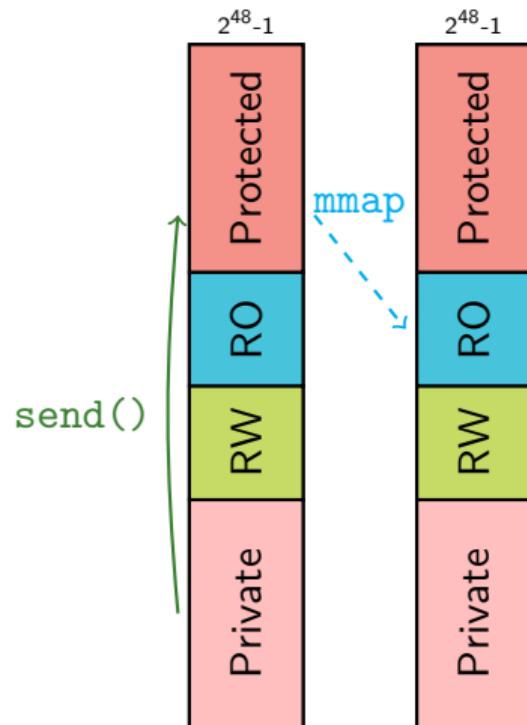
Polyhedral Analysis + GEMM: ☺

- Range checks using polyhedral analysis
 - Techniques from auto parallelization
 - e.g. Given `base`, conclude `base[tid]`
- **Minimal** overheads for ML workloads
 - Mostly regular
- Complex programs: add runtime checks
 - e.g. indirect heap references `a[b[i]]`
- Impl challenge: complexity of analysing directly on binaries



Efficient IPC: secure *direct memcpy*

- Useful primitive for multi-stage pipelines
 - Components from multiple vendors
- Validation enforces proper IPC region
 - Trusted primitive `send()`
 - Sender's **Protected**
 - `mmap` to receivers's **RO**
 - Avoid double encryption/decryption across the boundary of enclaves

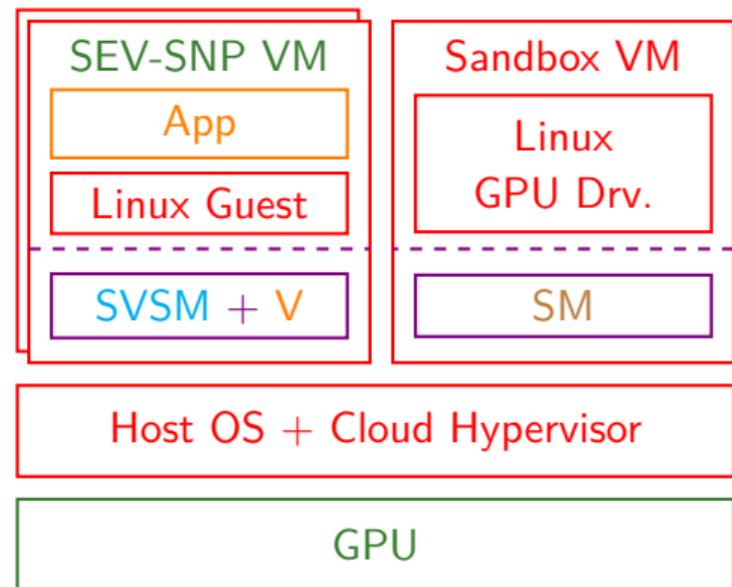


Agenda

- 1 Introduction
- 2 Assumptions & Background
- 3 Design & Implementation
- 4 Evaluation & Experience
- 5 Conclusion

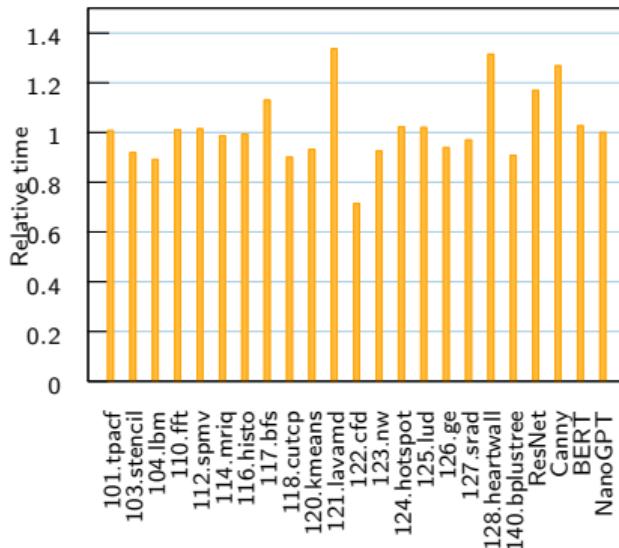
Security Monitors minimize TCB

- 2 EPYC 7433 CPUs, 1 AMD RX6900XT GPU
- Linux 5.17, ROCm 5.4.0
- TCB of Honeycomb (~82 KLOC): 18x smaller
 - Linux kernel ~1.5 MLOC
 - Core functionalities
 - Drivers (AMDGPU) and libraries (DRM & TTM)
 - Userspace runtime (ROCM) ~400 KLOC



Static validation is efficient

- 5 benchmark suites, HPC, CV, ML (DNN/Transformer). 23 apps in total
- Relative perf from 0.71-1.31 compared to Linux stack
 - breakdowns in paper
- Efficient on ML workloads
 - 2% overheads for BERT / NanoGPT
- Spent most time on GEMM kernels
 - polyhedral analysis works well
- Modest dev. effort to pass validations



Conclusion

- Honeycomb supports secure and efficient GPU executions
- Static analysis (Validation) is a practical and flexible technique for GPU apps
 - Honeycomb enhances security via co-designing validation + OS support
 - Efficient on real-world workloads
- The end-to-end SW/HW stack for GPU evolves quickly
 - A promising technique to explore novel designs